

Case Study 2 : Customer Transaction

To-do

Case Study Description

Let us take up the CUSTOMER and TRANSACTIONS table we have created in the Let's Do Together section. Let us solve the following use cases using these tables :-

1. Find out the number of transaction done by each customer (These should be take up in module 8 itself)

We have started hive shell by using command : **hive**

```
[acadgild@localhost ~]$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/hive-common-2.3.2.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive>
```

Then we have used custom database to create tables by using command: **use custom**

```
hive> use custom;
OK
```

Then we have created table CUSTOMER as shown below :

```
hive> CREATE TABLE CUSTOMER(
>   custid INT,
>   fname STRING,
>   lname STRING,
>   age INT,
>   profession STRING
> )
> row format delimited fields terminated by ',';
OK
Time taken: 0.332 seconds
```

Case Study 2 : Customer Transaction

We could verify that CUSTOMER table has been created successfully by using command :
SHOW TABLES

```
hive> show tables;  
OK  
array_concat  
array_demo  
array_demo2  
college  
customer  
olympics  
Time taken: 0.121 seconds, Fetched: 6 row(s)
```

Then we have inserted data into customer table from local file custs.txt as shown below :

LOAD DATA LOCAL INPATH '/home/acadgild/hive/custs.txt' into table CUSTOMER;

```
hive> LOAD DATA LOCAL INPATH '/home/acadgild/hive/custs.txt' into table CUSTOMER;  
Loading data to table custom.customer  
OK  
Time taken: 3.282 seconds
```

After inserting data into **customer** table, we are verifying that data is present in customer table by fetching rows by using query :

select * from customer;

You could see all 10 records in customer table.

Before that, we have set column header to TRUE so that we can have column headers along with output by using command :

set hive.cli.print.header = TRUE;

```
hive> set hive.cli.print.header = TRUE;  
hive> select * from customer;  
OK  
customer.custid customer.fname customer.lname customer.age customer.profession  
4000001 Kristina Chung 55 Pilot  
4000002 Paige Chen 74 Teacher  
4000003 Sherri Melton 34 Firefighter  
4000004 Gretchen Hill 66 Computer hardware engineer  
4000005 Karen Puckett 74 Lawyer  
4000006 Patrick Song 42 Veterinarian  
4000007 Elsie Hamilton 43 Pilot  
4000008 Hazel Bender 63 Carpenter  
4000009 Malcolm Wagner 39 Artist  
4000010 Dolores McLaughlin 60 Writer  
Time taken: 0.612 seconds, Fetched: 10 row(s)
```

Case Study 2 : Customer Transaction

Similarly we have created TRANSACTIONS table by using below query :

```
hive> CREATE TABLE TRANSACTIONS (txnno INT, txndate STRING, custno INT, amount DOUBLE, category STRING, product STRING, city STRING, state STRING, spendby STRING)
> row format delimited fields terminated by ',';
OK
Time taken: 0.755 seconds
```

Then we have inserted data into **transactions** table from local file txns.txt as shown below :

LOAD DATA LOCAL INPATH '/home/acadgild/hive/txns.txt' into table TRANSACTIONS;

```
hive> LOAD DATA LOCAL INPATH '/home/acadgild/hive/txns.txt' into table TRANSACTIONS;
Loading data to table custom.transactions
OK
```

After inserting data into **transactions** table, we are verifying that data is present in **transactions** table by fetching rows by using query :

select * from transactions;

```
hive> select * from TRANSACTIONS;
OK
transactions.txnno    transactions.txndate  transactions.custno  transactions.amount  transactions.category  transactions.prod
uct    transactions.city    transactions.state    transactions.spendby
0      06-26-2011          4000001 40.33 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee credit
1      05-26-2011          4000002 198.44 Exercise & Fitness Weightlifting Gloves Long Beach California credit
2      06-01-2011          4000002 5.58 Exercise & Fitness Weightlifting Machine Accessories Anaheim California credit
3      06-05-2011          4000003 198.19 Gymnastics Gymnastics Rings Milwaukee Wisconsin credit
4      12-17-2011          4000002 98.81 Team Sports Field Hockey Nashville Tennessee credit
5      02-14-2011          4000004 193.63 Outdoor Recreation Camping & Backpacking & Hiking Chicago Illinois credit
6      10-28-2011          4000005 27.89 Puzzles Jigsaw Puzzles Charleston South Carolina credit
7      07-14-2011          4000006 96.01 Outdoor Play Equipment Sandboxes Columbus Ohio credit
8      01-17-2011          4000006 10.44 Winter Sports Snowmobiling Des Moines Iowa credit
9      05-17-2011          4000006 152.46 Jumping Bungee Jumping St. Petersburg Florida credit
10     05-29-2011          4000007 180.28 Outdoor Recreation Archery Reno Nevada credit
11     06-18-2011          4000009 121.39 Outdoor Play Equipment Swing Sets Columbus Ohio credit
12     02-08-2011          4000009 41.52 Indoor Games Bowling San Francisco California credit
13     03-13-2011          4000010 107.8 Team Sports Field Hockey Honolulu Hawaii credit
14     02-25-2011          4000010 36.81 Gymnastics Vaulting Horses Los Angeles California credit
15     10-20-2011          4000001 137.64 Combat Sports Fencing Honolulu Hawaii credit
16     05-28-2011          4000010 35.56 Exercise & Fitness Free Weight Bars Columbia South Carolina credit
17     10-18-2011          4000008 75.55 Water Sports Scuba Diving & Snorkeling Omaha Nebraska credit
18     11-18-2011          4000008 88.65 Team Sports Baseball Salt Lake City Utah credit
19     08-28-2011          4000008 51.81 Water Sports Life Jackets Newark New Jersey credit
20     06-29-2011          4000005 41.55 Exercise & Fitness Weightlifting Belts New Orleans Louisiana credit
21     02-14-2011          4000005 45.79 Air Sports Parachutes New York New York credit
22     10-10-2011          4000009 19.64 Water Sports Kitesurfing Saint Paul Minnesota credit
23     05-02-2011          4000009 99.5 Gymnastics Gymnastics Rings Springfield Illinois credit
24     06-10-2011          4000003 151.2 Water Sports Surfing Plano Texas credit
25     10-14-2011          4000009 144.2 Indoor Games Darts Phoenix Arizona credit
26     10-11-2011          4000009 31.58 Combat Sports Wrestling Orange California credit
27     09-29-2011          4000010 66.4 Games Mahjong Fremont California credit
28     05-12-2011          4000008 79.78 Team Sports Cricket Lexington Kentucky credit
29     06-03-2011          4000001 126.9 Outdoor Recreation Hunting Phoenix Arizona credit
30     03-14-2011          4000001 47.05 Water Sports Swimming Lincoln Nebraska credit
31     11-28-2011          4000008 5.03 Games Dice & Dice Sets Los Angeles California credit
```

Then to find out the number of transaction done by each customer, we have used below query :
We have used JOIN with customer and transactions tables and grouped it by fname in customer.
Then we have fetched fname and count from this JOIN query.

select a.fname first_name, count(a.fname) count from CUSTOMER a join TRANSACTIONS b on a.custid =b.custno group by a.fname;

Case Study 2 : Customer Transaction

```
hive> select a.fname first_name, count(a.fname) count from CUSTOMER a join TRANSACTIONS b on a.custid =b.custno group by a.fname;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180818181554_1e5a71c8-8be9-424e-a7ee-0383b679b3e0
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/Static
LoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/i
mpl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-08-18 18:16:12 Starting to launch local task to process map join; maximum memory = 518979584
2018-08-18 18:16:16 Dump the side-table for tag: 0 with group count: 10 into file: file:/tmp/acadgild/348b012f-5364-4be0-b3c2-3997bea
f25e2/hive_2018-08-18_18-15-54_595_634880076181857495-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile30--.hashtable
2018-08-18 18:16:16 Uploaded 1 File to: file:/tmp/acadgild/348b012f-5364-4be0-b3c2-3997beaf25e2/hive_2018-08-18_18-15-54_595_63488007
6181857495-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile30--.hashtable (556 bytes)
2018-08-18 18:16:16 End of local task; Time Taken: 4.68 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1534569566276_0005, Tracking URL = http://localhost:8088/proxy/application_1534569566276_0005/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1534569566276_0005
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-08-18 18:16:35,772 Stage-2 map = 0%, reduce = 0%
2018-08-18 18:16:51,120 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.25 sec
2018-08-18 18:17:06,218 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 7.83 sec
MapReduce Total cumulative CPU time: 7 seconds 830 msec
Ended Job = job_1534569566276_0005
```

So the Final output is :

first_name	count
------------	-------

Dolores	6
Elsie	6
Gretchen	5
Hazel	10
Karen	5
Kristina	8
Malcolm	6
Paige	6
Patrick	5
Sherri	3

```
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 7.83 sec HDFS Read: 17689 HDFS Write: 301 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 830 msec
OK
first_name      count
Dolores 6
Elsie 6
Gretchen 5
Hazel 10
Karen 5
Kristina 8
Malcolm 6
Paige 6
Patrick 5
Sherri 3
Time taken: 72.818 seconds, Fetched: 10 row(s)
```

Case Study 2 : Customer Transaction

2. Create a new table called TRANSACTIONS_COUNT. This table should have

3 fields - custid, fname and count. (Again to be done in module 8)

We have created TRANSACTIONS_COUNT table with three columns : custid, fname and count by using below query :

```
hive> CREATE TABLE TRANSACTIONS_COUNT(
> custid INT,
> fname STRING,
> count INT
> )
> row format delimited fields terminated by ',';
OK
Time taken: 0.459 seconds
```

3. Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above. (This has to be done in module 9).

We have inserted data into TRANSACTIONS_COUNT table by using similar select query which we have used for query in question 1.

insert overwrite table TRANSACTIONS_COUNT

select a.custid CUSTID, a.fname FNAME, count(a.fname) COUNT from CUSTOMER a JOIN TRANSACTIONS b on a.custid =b.custno group by a.fname,a.custid;

```
hive> insert overwrite table TRANSACTIONS_COUNT
> select a.custid CUSTID, a.fname FNAME, count(a.fname) COUNT from CUSTOMER a JOIN TRANSACTIONS b on a.custid =b.custno group by a.fname,a.custid;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180818193727_415ba556-7d09-441b-be37-009d8dda3f67
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-08-18 19:37:52 Starting to launch local task to process map join; maximum memory = 518979584
2018-08-18 19:37:57 Dump the side-table for tag: 0 with group count: 10 into file: file:/tmp/acadgild/f62a0f9d-db32-4b1b-bfed-3d033af899a7/hive_2018-08-18_19-37-27_484_1060530934380108427-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile00--.hashtable
2018-08-18 19:37:57 Uploaded 1 File to: file:/tmp/acadgild/f62a0f9d-db32-4b1b-bfed-3d033af899a7/hive_2018-08-18_19-37-27_484_1060530934380108427-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile00--.hashtable (556 bytes)
2018-08-18 19:37:57 End of local task; Time Taken: 5.628 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1534569566276_0006, Tracking URL = http://localhost:8088/proxy/application_1534569566276_0006/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1534569566276_0006
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-08-18 19:38:33,703 Stage-2 map = 0%, reduce = 0%
2018-08-18 19:39:04,615 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 6.22 sec
```

Case Study 2 : Customer Transaction

```
2018-08-18 10:39:25.457 Stage-2 map = 100%, reduce = 73%, Cumulative CPU 11.26 sec
2018-08-18 10:39:26.507 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 11.85 sec
MapReduce Total cumulative CPU time: 11 seconds 850 msec
Ended Job = job_1534569566276_0006
Loading data to table custom.transactions_count
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 11.85 sec HDFS Read: 18814 HDFS Write: 256 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 850 msec
OK
Time taken: 124.365 seconds
```

After inserting data into TRANSACTIONS_COUNT table, we are verifying that data is present in TRANSACTIONS_COUNT table by fetching rows by using select query :

```
select * from TRANSACTIONS_COUNT;
```

You could see all 10 records in customer table.

```
hive> select * from TRANSACTIONS_COUNT;
OK
transactions_count.custid      transactions_count.fname      transactions_count.count
40000001 Kristina            8
40000002 Paige               6
40000003 Sherri              3
40000004 Gretchen            5
40000005 Karen               5
40000006 Patrick             5
40000007 Elsie               6
40000008 Hazel               10
40000009 Malcolm             6
40000010 Dolores             6
Time taken: 0.521 seconds, Fetched: 10 row(s)
```

4. Now lets make the TRANSACTIONS_COUNT table Hbase complaint. In the sense, use Ser Des And Storate handler features of hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase. (This has to be done in module 10)

So we have created a table TRANSACTIONS_COUNT2 similar to TRANSACTIONS_COUNT as HBase compliant as shown below :

```
CREATE TABLE TRANSACTIONS_COUNT2 (custid int, fname string, count int)
  STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
  WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,customer_details:fname,customer_details:count")
  TBLPROPERTIES ("hbase.table.name" = "TRANSACTIONS");
```

```
hive> CREATE TABLE TRANSACTIONS_COUNT2(custid int, fname string, count int)
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,customer_details:fname,customer_details:count")
> TBLPROPERTIES ("hbase.table.name" = "TRANSACTIONS");
OK
Time taken: 2.899 seconds
hive> select * from TRANSACTIONS_COUNT2;
OK
transactions_count2.custid      transactions_count2.fname      transactions_count2.count
Time taken: 0.709 seconds
hive> describe TRANSACTIONS_COUNT2;
OK
col_name      data_type      comment
custid        int
fname         string
count         int
Time taken: 0.265 seconds, Fetched: 3 row(s)
```

Case Study 2 : Customer Transaction

You could see that we have all three columns : custid, fname,count present in TRANSACTIONS_COUNT2 table as in TRANSACTIONS_COUNT.

Here we have used **STORED BY** to make this Hive table as HBase compliant.

By using **TBLPROPERTIES**, we could create table TRANSCATIONS in HBase and by using **SERDEPROPERTIES**, we could map columns of TRANSACTIONS_COUNT2 table in Hive with columns of TRANSACTIONS in HBase.

Then we have connected to Hbase shell and verified that TRANSACTIONS table in HBase has been created by using commands :

scan 'TRANSACTIONS' which shows that table has been created successfully and

describe 'TRANSACTIONS' which shows column family is customer_details and other details.

```
hbase(main):012:0> scan 'TRANSACTIONS'
ROW
0 row(s) in 0.0340 seconds

hbase(main):013:0> describe 'TRANSACTIONS'
Table TRANSACTIONS is ENABLED
TRANSACTIONS
COLUMN FAMILIES DESCRIPTION
{NAME => 'customer_details', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
1 row(s) in 0.1080 seconds
```

5. Now insert the data in TRANSACTIONS_COUNT table using the query in step

3 again, this should populate the Hbase TRANSACTIONS table automatically

(This has to be done in module 10)

Then to insert data into HBase table TRANSACTIONS, we have loaded data into TRANSACTIONS_COUNT2 from TRANSACTIONS_COUNT as

insert overwrite table transactions_count2 select * from transactions_count;

Case Study 2 : Customer Transaction

```
hive> insert overwrite table TRANSACTIONS_COUNT2 select * from TRANSACTIONS_COUNT;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180819113914_6533aa61-e33c-4c15-9f31-d790c06e0c5f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1534655884402_0002, Tracking URL = http://localhost:8088/proxy/application_1534655884402_0002/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1534655884402_0002
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2018-08-19 11:39:50,142 Stage-3 map = 0%, reduce = 0%
2018-08-19 11:40:14,212 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 4.56 sec
MapReduce Total cumulative CPU time: 5 seconds 450 msec
Ended Job = job_1534655884402_0002
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 5.45 sec HDFS Read: 4894 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 450 msec
OK
transactions_count.custid      transactions_count.fname      transactions_count.count
Time taken: 63.384 seconds
hive>
```

We could verify that data has been inserted successfully into table TRANSACTIONS_COUNT2 by using below select query :

select * from TRANSACTIONS_COUNT2;

```
hive> select * from TRANSACTIONS_COUNT2;
OK
transactions_count2.custid      transactions_count2.fname      transactions_count2.count
4000001 Kristina                8
4000002 Paige                  6
4000003 Sherri                 3
4000004 Gretchen               5
4000005 Karen                  5
4000006 Patrick                5
4000007 Elsie                  6
4000008 Hazel                  10
4000009 Malcolm                6
4000010 Dolores                6
Time taken: 0.989 seconds, Fetched: 10 row(s)
```

After Inserting data into TRANSACTIONS_COUNT2, we have verified that same data has been inserted into TRANSACTIONS table in HBase by using command :

scan 'TRANSACTIONS'

We could see that all 10 rows have been inserted into TRANSACTIONS table successfully.

```
hbase(main):014:0> scan 'TRANSACTIONS'
ROW
4000001      column=customer_details:count, timestamp=1534659014582, value=8
4000001      column=customer_details:fname, timestamp=1534659014582, value=Kristina
4000002      column=customer_details:count, timestamp=1534659014582, value=6
4000002      column=customer_details:fname, timestamp=1534659014582, value=Paige
4000003      column=customer_details:count, timestamp=1534659014582, value=3
4000003      column=customer_details:fname, timestamp=1534659014582, value=Sherri
4000004      column=customer_details:count, timestamp=1534659014582, value=5
4000004      column=customer_details:fname, timestamp=1534659014582, value=Gretchen
4000005      column=customer_details:count, timestamp=1534659014582, value=5
4000005      column=customer_details:fname, timestamp=1534659014582, value=Karen
4000006      column=customer_details:count, timestamp=1534659014582, value=5
4000006      column=customer_details:fname, timestamp=1534659014582, value=Patrick
4000007      column=customer_details:count, timestamp=1534659014582, value=6
4000007      column=customer_details:fname, timestamp=1534659014582, value=Elsie
4000008      column=customer_details:count, timestamp=1534659014582, value=10
4000008      column=customer_details:fname, timestamp=1534659014582, value=Hazel
4000009      column=customer_details:count, timestamp=1534659014582, value=6
4000009      column=customer_details:fname, timestamp=1534659014582, value=Malcolm
4000010      column=customer_details:count, timestamp=1534659014582, value=6
4000010      column=customer_details:fname, timestamp=1534659014582, value=Dolores
10 row(s) in 0.1750 seconds
```


Case Study 2 : Customer Transaction

6. Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level.

We have written two JAVA API codes :

1. **ScanTableAdvanced.java** : This will scan entire HBase table TRANSACTIONS.
2. **GetOperations.java** : This will fetch data in TRANSACTIONS table for a particular rowkey.

1. Below is the java code of **ScanTableAdvanced** :

```
import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.client.Result;
import org.apache.hadoop.hbase.client.ResultScanner;
import org.apache.hadoop.hbase.client.Scan;

public class ScanTableAdvanced {
    public static void main(String[] args) throws IOException,
        InterruptedException {
        Configuration conf = HBaseConfiguration.create();

        System.out.println("Creating HTable instance to 'TRANSACTIONS'...");
        HTable table = new HTable(conf, "TRANSACTIONS");

        System.out.println("Creating scan object to scan TRANSACTIONS
table...");
        Scan scan = new Scan();

        System.out.println("Scanner Caching at table level: " +
table.getScannerCaching());
        scan.setCaching(1);
        scan.setBatch(2);

        System.out.println("Scanner Caching at scan object level: " +
scan.getCaching());

        System.out.println("Getting a result scanner object...");
        ResultScanner rs = table.getScanner(scan);

        for (Result r : rs) {
            System.out.println("Result: " + r);
        }

        System.out.println("Closing Scanner instance...");
        rs.close();
    }
}
```

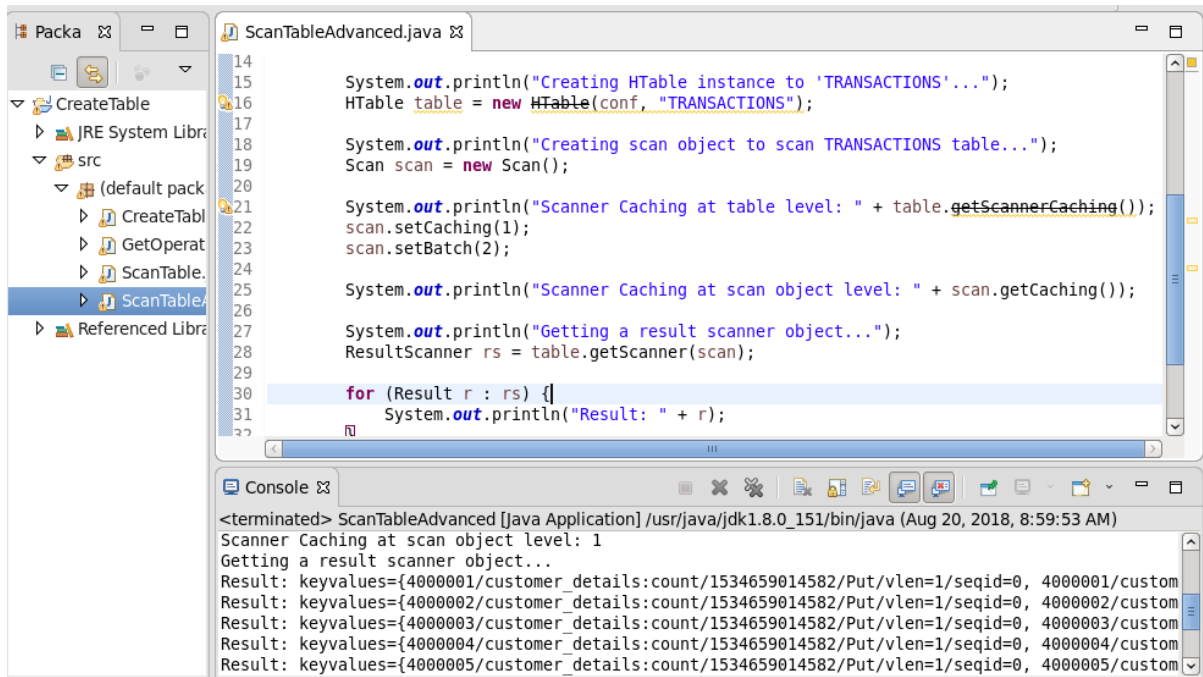
Case Study 2 : Customer Transaction

Below is the output at Console :

```
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/spark/spark-
2.2.1-bin-hadoop2.7/jars/slf4j-log4j12-
1.7.16.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-
1.2.6/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
log4j:WARN No appenders could be found for logger
(org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig
for more info.
Creating HTable instance to 'TRANSACTIONS'...
Creating scan object to scan TRANSACTIONS table...
Scanner Caching at table level: 2147483647
Scanner Caching at scan object level: 1
Getting a result scanner object...
Result:
keyvalues={4000001/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000001/customer_details:fname/1534659014582/Put/vlen=8/seqid=0}
Result:
keyvalues={4000002/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000002/customer_details:fname/1534659014582/Put/vlen=5/seqid=0}
Result:
keyvalues={4000003/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000003/customer_details:fname/1534659014582/Put/vlen=6/seqid=0}
Result:
keyvalues={4000004/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000004/customer_details:fname/1534659014582/Put/vlen=8/seqid=0}
Result:
keyvalues={4000005/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000005/customer_details:fname/1534659014582/Put/vlen=5/seqid=0}
Result:
keyvalues={4000006/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000006/customer_details:fname/1534659014582/Put/vlen=7/seqid=0}
Result:
keyvalues={4000007/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000007/customer_details:fname/1534659014582/Put/vlen=5/seqid=0}
Result:
keyvalues={4000008/customer_details:count/1534659014582/Put/vlen=2/seqi
d=0, 4000008/customer_details:fname/1534659014582/Put/vlen=5/seqid=0}
Result:
keyvalues={4000009/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000009/customer_details:fname/1534659014582/Put/vlen=7/seqid=0}
Result:
keyvalues={4000010/customer_details:count/1534659014582/Put/vlen=1/seqi
d=0, 4000010/customer_details:fname/1534659014582/Put/vlen=7/seqid=0}
Closing Scanner instance...
```

Case Study 2 : Customer Transaction

Below is the screenshot of ScanTableAdvanced.java code and it's console :



The screenshot shows an IDE with a project named 'Packa'. The 'src' folder contains files: 'CreateTable', 'GetOperat', 'ScanTable', and 'ScanTableAdvanced.java'. The 'ScanTableAdvanced.java' file is open, showing the following code:

```
14
15 System.out.println("Creating HTable instance to 'TRANSACTIONS'...");
16 HTable table = new HTable(conf, "TRANSACTIONS");
17
18 System.out.println("Creating scan object to scan TRANSACTIONS table...");
19 Scan scan = new Scan();
20
21 System.out.println("Scanner Caching at table level: " + table.getScannerCaching());
22 scan.setCaching(1);
23 scan.setBatch(2);
24
25 System.out.println("Scanner Caching at scan object level: " + scan.getCaching());
26
27 System.out.println("Getting a result scanner object...");
28 ResultScanner rs = table.getScanner(scan);
29
30 for (Result r : rs) {
31     System.out.println("Result: " + r);
32 }
```

The console output shows the execution of the code:

```
<terminated> ScanTableAdvanced [Java Application] /usr/java/jdk1.8.0_151/bin/java (Aug 20, 2018, 8:59:53 AM)
Scanner Caching at scan object level: 1
Getting a result scanner object...
Result: keyvalues={4000001/customer_details:count/1534659014582/Put/vlen=1/seqid=0, 4000001/custom
Result: keyvalues={4000002/customer_details:count/1534659014582/Put/vlen=1/seqid=0, 4000002/custom
Result: keyvalues={4000003/customer_details:count/1534659014582/Put/vlen=1/seqid=0, 4000003/custom
Result: keyvalues={4000004/customer_details:count/1534659014582/Put/vlen=1/seqid=0, 4000004/custom
Result: keyvalues={4000005/customer_details:count/1534659014582/Put/vlen=1/seqid=0, 4000005/custom
```

This above java code will scan entire TRANSACTIONS table. Hence you could see in the output that all 10 rowkeys from 4000001 to 4000010 with their respective count and fname column values are present.

Case Study 2 : Customer Transaction

2. Below is the java code of **getOperations** :

```
import java.io.IOException;
import java.util.Map;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.KeyValue;
import org.apache.hadoop.hbase.client.Get;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.util.Bytes;
import org.apache.hadoop.hbase.client.Result;

public class GetOperations {

    public static void main(String[] args) throws IOException {

        Configuration conf = HBaseConfiguration.create();

        System.out.println("Creating HTable instance to 'TRANSACTIONS'...");
        HTable table = new HTable(conf, "TRANSACTIONS");

        System.out.println("Creating get object to retrieve the row with key
p1...");
        Get get = new Get(Bytes.toBytes("4000008"));
        System.out.println("Default Time Range of get object: " +
get.getTimeRange());
        System.out.println("Will Block Cache be scanned: " +
get.getCacheBlocks());
        System.out.println("Default versions to fetch: " +
get.getMaxVersions());

        System.out.println("Setting versions to fetch to be 3");
        get.setMaxVersions(3);
        System.out.println("Versions to fetch: " + get.getMaxVersions());

        Result result = table.get(get);
        System.out.println("Result fetched: " + result);

        System.out.println("Fetching the most recent givenName...");
        String givenName =
Bytes.toString(result.getValue(Bytes.toBytes("customer_details"),
Bytes.toBytes("fname")));
        System.out.println("Given Name retrieved: " + givenName);

        System.out.println("Scanning across all the values of column
personal:givenName...");

        for (KeyValue kv :
result.getColumn(Bytes.toBytes("customer_details"), Bytes.toBytes("fname"))) {
            System.out.println("Timestamp: " + kv.getTimestamp() + "\t" +
"Value of giveName: " + Bytes.toString(kv.getValue()));
        }

        System.out.println("Scanning across all the values of key p1...");
    }
}
```

Case Study 2 : Customer Transaction

```
        for (KeyValue kv : result.raw()) {
            System.out.println("Row Key: " + Bytes.toString(kv.getRow()) +
                ", Column Family: " + Bytes.toString(kv.getFamily()) +
                ", Column Name: " +
                Bytes.toString(kv.getQualifier()) + ", Value: " + Bytes.toString(kv.getValue()) +
                ", Timestamp: " + kv.getTimestamp() );
        }

        Map resultMap = result.getMap();
        for (Object key : resultMap.keySet()) {
            System.out.println("Key: " + Bytes.toString((byte[]) key));
        }

        System.out.println(resultMap.get(Bytes.toBytes("TRANSACTIONS")));
        System.out.println(resultMap.values());
        System.out.println(result.getNoVersionMap());
    }
}
```

Below is the screenshot of getOperations.java code and it's console :

```
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/spark/spark-2.2.1-bin-hadoop2.7/jars/slf4j-log4j12-1.7.16.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
log4j:WARN No appenders could be found for logger
(org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig
for more info.
Creating HTable instance to 'TRANSACTIONS'...
Creating get object to retrieve the row with key p1...
Default Time Range of get object: maxStamp=9223372036854775807,
minStamp=0
Will Block Cache be scanned: true
Default versions to fetch: 1
Setting versions to fetch to be 3
Versions to fetch: 3
Result fetched:
keyvalues={4000008/customer_details:count/1534659014582/Put/vlen=2/seqid=0, 4000008/customer_details:fname/1534659014582/Put/vlen=5/seqid=0}
Fetching the most recent givenName...
Given Name retrieved: Hazel
Scanning across all the values of column personal:givenName...
Timestamp: 1534659014582      Value of giveName: Hazel
```

Case Study 2 : Customer Transaction

Scanning across all the values of key pl...

Row Key: 4000008, Column Family: customer_details, Column Name: count, Value: 10, Timestamp: 1534659014582

Row Key: 4000008, Column Family: customer_details, Column Name: fname, Value: Hazel, Timestamp: 1534659014582

Key: customer_details

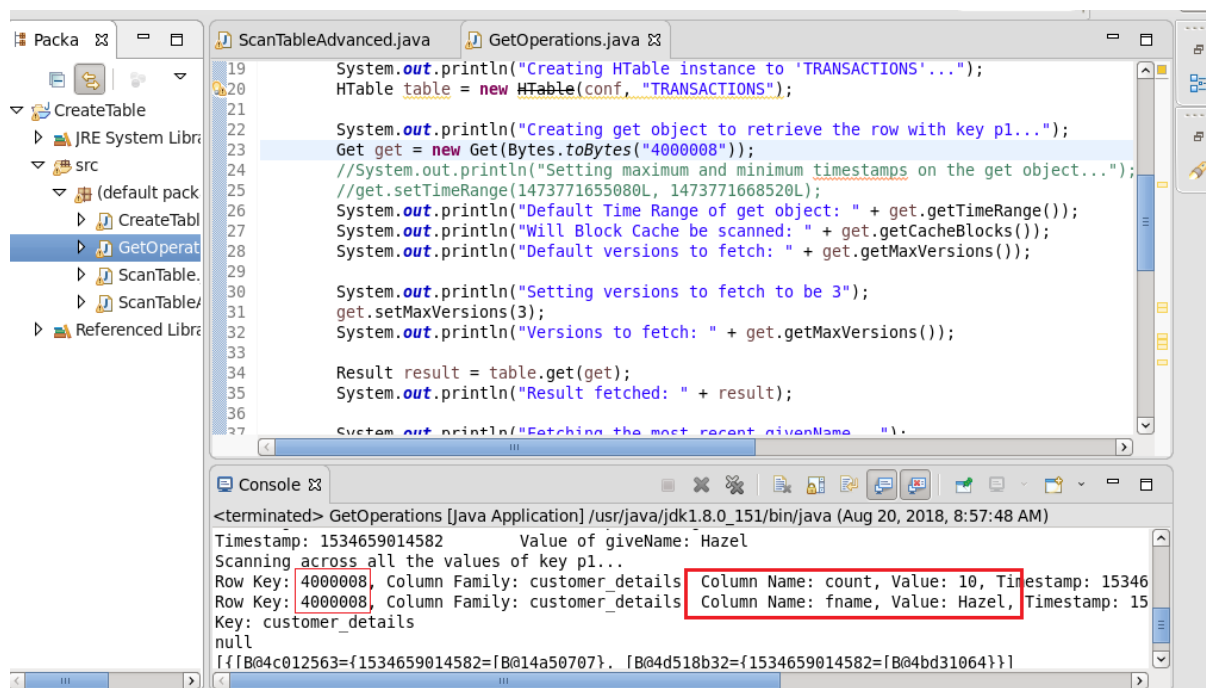
null

```
{ [B@4c012563={1534659014582=[B@14a50707},
```

```
[B@4d518b32={1534659014582=[B@4bd31064} }]
```

```
{ [B@e3c0e40={ [B@4c012563=[B@14a50707, [B@4d518b32=[B@4bd31064} }
```

Below is the screenshot of GetOperations.java code and it's console :



```
19 System.out.println("Creating HTable instance to 'TRANSACTIONS'...");
20 HTable table = new HTable(conf, "TRANSACTIONS");
21
22 System.out.println("Creating get object to retrieve the row with key pl...");
23 Get get = new Get(Bytes.toBytes("4000008"));
24 //System.out.println("Setting maximum and minimum timestamps on the get object...");
25 //get.setTimeRange(1473771655080L, 1473771668520L);
26 System.out.println("Default Time Range of get object: " + get.getTimeRange());
27 System.out.println("Will Block Cache be scanned: " + get.getCacheBlocks());
28 System.out.println("Default versions to fetch: " + get.getMaxVersions());
29
30 System.out.println("Setting versions to fetch to be 3");
31 get.setMaxVersions(3);
32 System.out.println("Versions to fetch: " + get.getMaxVersions());
33
34 Result result = table.get(get);
35 System.out.println("Result fetched: " + result);
36
37 System.out.println("Fetching the most recent givenName: ");
```

```
<terminated> GetOperations [Java Application] /usr/java/jdk1.8.0_151/bin/java (Aug 20, 2018, 8:57:48 AM)
Timestamp: 1534659014582      Value of giveName: Hazel
Scanning across all the values of key pl...
Row Key: 4000008, Column Family: customer_details, Column Name: count, Value: 10, Timestamp: 15346
Row Key: 4000008, Column Family: customer_details, Column Name: fname, Value: Hazel, Timestamp: 15
Key: customer_details
null
[[B@4c012563={1534659014582=[B@14a50707}, [B@4d518b32={1534659014582=[B@4bd31064}]]]
```

This above java code will fetch TRANSACTIONS table only having rowkey **4000008** as we have mentioned in our code to fetch data with rowkey as 4000008. Hence you could see in the output that count value is 10 and fname value is Hazel.